

This file will give a brief tutorial on how to install and use the ONAM package.

Install package:

```
install.packages("devtools")  
library(devtools)  
devtools::install_github("Koehlibert/ONAM")
```

If this is the first time using Keras or TensorFlow in R, you need to run `keras::install_keras()`.

The following packages are required for data examples, visualization and machine learning models used for demonstration:

```
library(ONAM)  
library(mlbench)  
library(MASS)  
library(gbm)  
library(e1071)  
library(caret)  
library(xgboost)  
library(dplyr)
```

We demonstrate the algorithm by explaining a gradient boosting machine for probability prediction of diabetes, and a xgboost model for prediction of housing prices using the boston housing dataset.

Model specification, evaluation and visualization will be demonstrated.

Binary classification gradient boosting machine for diabetes prediction

The Pima Indians Diabetes Database is a dataset from the Indian National Institute of Diabetes and Digestive and Kidney Diseases. <https://www.mdpi.com/2076-3417/9/21/4604>

```
data(PimaIndiansDiabetes)
diabetes_data <- PimaIndiansDiabetes %>%
  mutate(diabetes = ifelse(diabetes == "pos", 1, 0)) %>%
  mutate(across(!diabetes, as.numeric)) %>%
  filter(mass > 0)

head(diabetes_data)
```

##	pregnant	glucose	pressure	triceps	insulin	mass	pedigree	age	diabetes
## 1	6	148	72	35	0	33.6	0.627	50	1
## 2	1	85	66	29	0	26.6	0.351	31	0
## 3	8	183	64	0	0	23.3	0.672	32	1
## 4	1	89	66	23	94	28.1	0.167	21	0
## 5	0	137	40	35	168	43.1	2.288	33	1
## 6	5	116	74	0	0	25.6	0.201	30	0

It consists of 757 instances of multiple medical features and the diabetes status of each observation. There are 266 diabetes cases and 491 controls in the dataset.

We demonstrate how to use ONAM to explain the prediction of a gradient boosting machine for prediction of the diabetes status.

First, we fit a gradient boosting machine for prediction of diabetes probability.

```
# Train the GBM model (binary classification)
gbm_model <- gbm(diabetes ~ .,
  data = diabetes_data,
  distribution = "bernoulli",
  n.trees = 500,
  interaction.depth = 3,
  shrinkage = 0.01,
  cv.folds = 5)
summary(gbm_model, plotit = F)
```

##		var	rel.inf
## glucose	glucose	43.659244	
## mass	mass	19.757377	
## age	age	14.334010	
## pedigree	pedigree	9.600033	
## pregnant	pregnant	4.656626	
## insulin	insulin	3.459595	
## pressure	pressure	3.043466	
## triceps	triceps	1.489649	

Model performance:

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction    0    1
##           0 469 128
##           1  22 138
##
##           Accuracy : 0.8018
##           95% CI : (0.7716, 0.8297)
##           No Information Rate : 0.6486
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.5216
##
##  Mcnemar's Test P-Value : < 2.2e-16
##
##           Sensitivity : 0.9552
##           Specificity : 0.5188
##           Pos Pred Value : 0.7856
##           Neg Pred Value : 0.8625
##           Prevalence : 0.6486
##           Detection Rate : 0.6196
##           Detection Prevalence : 0.7886
##           Balanced Accuracy : 0.7370
##
##           'Positive' Class : 0
##
```

Fit ONAM model on gbm data

We explain the svm using the orthogonal additive model framework for binary classification:

```
Diabetes_formula <- diabetes ~ mod(glucose) + mod(insulin) +
  mod(pregnant) + mod(pressure) + mod(triceps) +
  mod(pedigree) + mod(mass) + mod(age) +
  mod(glucose, insulin) + mod(age, mass) + mod(.)
list_of_mods_Diabetes <- list(mod = ONAM::get_submodel)
gbm_expl <- onam(
  Diabetes_formula,
  list_of_deep_models_Diabetes,
  diabetes_data,
  gbm_model,
  prediction_function = function(model, data) {
    predict(model, data, n.trees = 500, type = "response")
  },
  target = "binary",
  n_ensemble = 20,
  progresstext = TRUE,
  epochs = 500,
  verbose = 0
)

#generate prediction object for saving
gbm_res <- predict(gbm_expl)
saveRDS(gbm_res, "gbm_res.RDS")

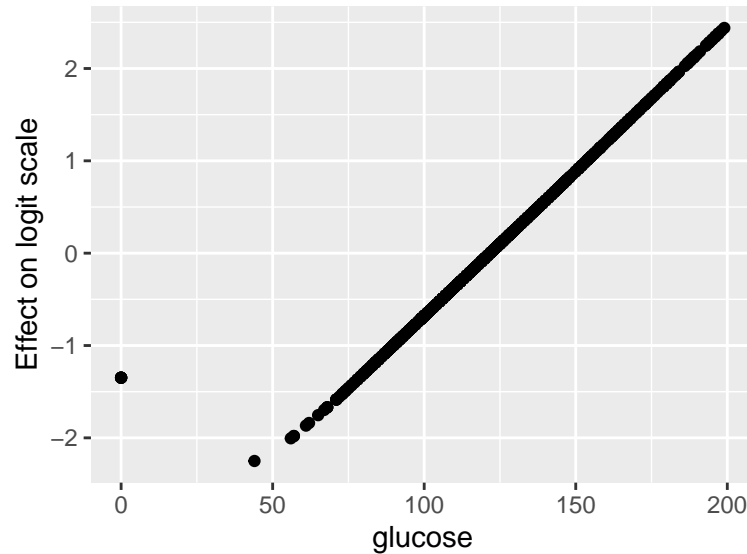
##
## Call:
## onam(formula = Diabetes_formula, list_of_deep_models = list_of_deep_models_Diabetes,
##      data = diabetes_data, model = gbm_model, prediction_function = function(model,
##      data) {
##        predict(model, data, n.trees = 500, type = "response")
##      }, target = "binary", n_ensemble = 10, epochs = 500, progresstext = TRUE,
##      verbose = 0)
##
## Correlation of onam probabilities with original model predicted probabilities: 0.9889
## Number of ensemble members: 10
## I_1: 0.9726; I_2: 0.0101
## Degree of interpretability: 0.9827
```

Example of effect plots

The shown effects are on the logit scale.

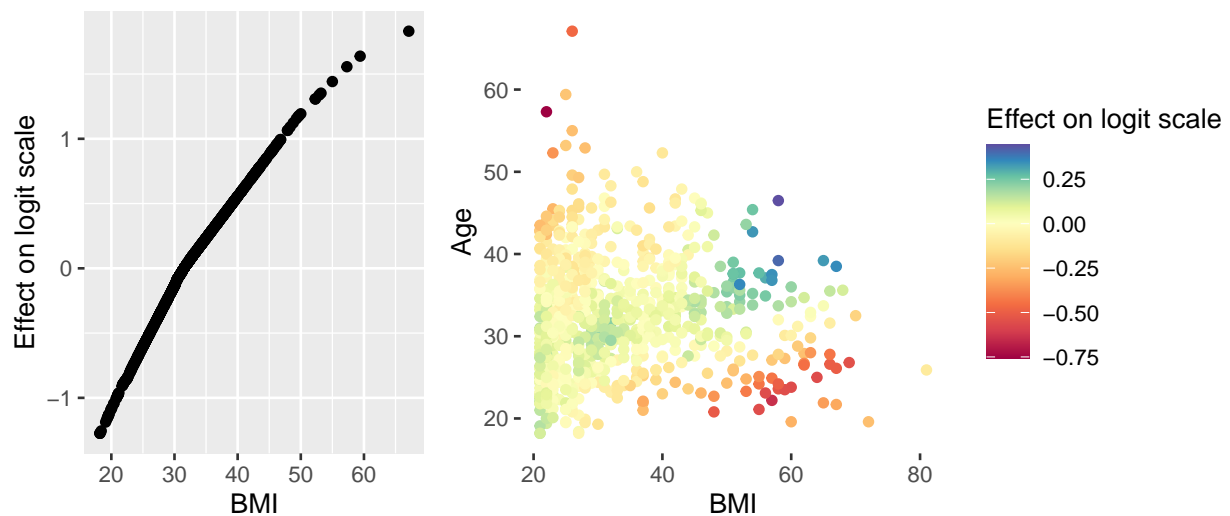
A higher plasma glucose concentration is associated with a higher risk of diabetes:

```
plot_main_effect(gbm_res, "glucose")
```



A higher BMI is associated with a higher risk of diabetes, but the increase is lower in younger people.

```
plot_main_effect(gbm_res, "mass") + xlab("BMI")  
  
plot_inter_effect(gbm_res, "age", "mass", interpolate = FALSE) +  
  xlab("BMI") + ylab("Age")
```



Prediction of Boston housing prices using xgboost

The Boston dataset contains data collected by the US Census Service on socioecological features of housing in Boston, MA. It contains 13 features along with the median value of owner-occupied homes, which will be the outcome in this section.

```
data("Boston")
```

```
head(Boston)
```

```
##      crim zn indus chas   nox    rm  age    dis rad tax ptratio  black lstat
## 1 0.00632 18  2.31    0 0.538 6.575 65.2 4.0900   1 296    15.3 396.90  4.98
## 2 0.02731  0  7.07    0 0.469 6.421 78.9 4.9671   2 242    17.8 396.90  9.14
## 3 0.02729  0  7.07    0 0.469 7.185 61.1 4.9671   2 242    17.8 392.83  4.03
## 4 0.03237  0  2.18    0 0.458 6.998 45.8 6.0622   3 222    18.7 394.63  2.94
## 5 0.06905  0  2.18    0 0.458 7.147 54.2 6.0622   3 222    18.7 396.90  5.33
## 6 0.02985  0  2.18    0 0.458 6.430 58.7 6.0622   3 222    18.7 394.12  5.21
##   medv
## 1 24.0
## 2 21.6
## 3 34.7
## 4 33.4
## 5 36.2
## 6 28.7
```

The dataset contains information 506 towns. We use xgboost to fit a prediction model for median housing prices per town.

```
train_Boston <- Boston %>% dplyr::select(-medv)
label_Boston <- Boston %>% dplyr::select(medv) %>% unlist()

xgb_train_Boston <- xgb.DMatrix(as.matrix(train_Boston),
                                label = label_Boston)

xgb_mod <-
  xgboost(
    xgb_train_Boston,
    nrounds = 10000,
    verbose = 0,
    objective = "reg:squarederror",
    eval_metric = "rmse"
  )
summary(xgb_mod)
```

```
##      Length Class      Mode
## handle          1 xgb.Booster.handle externalptr
## raw            7318563 -none-      raw
## niter           1 -none-      numeric
## evaluation_log    2 data.table    list
## call             15 -none-      call
## params            3 -none-      list
## callbacks         1 -none-      list
## feature_names     13 -none-      character
## nfeatures         1 -none-      numeric
```

```
preds <- predict(xgb_mod, xgb_train_Boston)
pricing_cor <- cor(preds, Boston$medv)
```

Correlation between housing prices and predictions: 1

Fit ONAM model on xgboost data

We explain the xgboost using the orthogonal additive model framework:

```
Boston_formula <-
  medv ~ mod(crim) + mod(zn) + mod(indus) + mod(chas) +
  mod(nox) + mod(rm) + mod(age) + mod(dis) + mod(rad) +
  mod(tax) + mod(ptratio) + mod(black) + mod(lstat) +
  mod(crim, dis) + mod(rm, age) + mod(zn, indus) +
  mod(.)

list_of_deep_models_Boston <- list(mod = ONAM::get_submodel)

categorical_features <- c("chas")

xgb_expl <-
  onam(
    Boston_formula,
    list_of_deep_models_Boston,
    Boston,
    # target = "binary",
    model = xgb_mod,
    model_data = xgb_train_Boston,
    categorical_features = categorical_features,
    n_ensemble = 20,
    epochs = 1000,
    progresstext = TRUE,
    verbose = 0
  )


xgb_res <- predict(xgb_expl)
saveRDS(xgb_res, "xgb_res.RDS")

##
## Call:
## onam(formula = Boston_formula, list_of_deep_models = list_of_deep_models_Boston,
##       data = Boston, model = xgb_mod, model_data = xgb_train_Boston,
##       categorical_features = categorical_features, n_ensemble = 20,
##       epochs = 1000, progresstext = TRUE, verbose = 0)
##
## Correlation of model prediction with outcome variable: 0.9942
## Number of ensemble members: 20
## I_1: 0.924; I_2: 0.0415
## Degree of interpretability: 0.9656
```

Examples of effect plots

Higher crime rate is associated with lower housing prices

```
plot_main_effect(xgb_res, "crim") + xlab("Crime rate (%)")
```



Lower number of rooms is associated with lower housing prices, but less so in (areas with) older houses.

```
plot_main_effect(xgb_res, "rm") + xlab("Av. number of rooms per dwelling")
plot_inter_effect(xgb_res, "rm", "age") +
  xlab("Av. number of rooms per dwelling") +
  ylab("% of units build prior to 1940")
```

